

---

# *Kubernetes* 企业级高可用部署

# 目录

目录.....	2
文档信息 .....	3
文档约定 .....	3
1、KUBERNETES 高可用项目介绍.....	3
2、项目架构设计 .....	3
2.1、项目主机信息 .....	3
2.2、项目架构图 .....	4
2.3、项目实施思路 .....	5
3、项目实施过程 .....	6
3.1、系统初始化 .....	6
3.2、配置部署 KEEPALIVED 服务 .....	7
3.3、配置部署 HAPROXY 服务 .....	10
3.4、配置部署 DOCKER 服务 .....	13
3.5、部署 KUBELET KUBEADM KUBECTL 工具 .....	13
3.6、部署 KUBERNETES MASTER .....	14
3.7、安装集群网络 .....	17
3.8、添加 MASTER 节点 .....	18
3.9、加入 KUBERNETES NODE.....	19
3.10、测试 KUBERNETES 集群.....	20
4、项目总结 .....	23

## 文档信息

文档作者	房佳亮
文档版本	Version1.0
文档版权	内部资料禁止传播
文档归类	Kubernetes 项目实战训练营
系统环境	CentOS-7.X-x86_64
作者邮箱	crushlinux@163.com
修订信息	2021-03-31
技术交流	

## 文档约定

[绿色背景]	知识重点
[红色背景]	错误警告
[黄色背景]	注意事项

### 执行命令

## 1、Kubernetes 高可用项目介绍

单 master 节点的可靠性不高，并不适合实际的生产环境。Kubernetes 高可用集群是保证 Master 节点中 API Server 服务的高可用。API Server 提供了 Kubernetes 各类资源对象增删改查的唯一访问入口，是整个 Kubernetes 系统的数据总线 and 数据中心。采用负载均衡（Load Balance）连接多个 Master 节点可以提供稳定容器云业务。

## 2、项目架构设计

### 2.1、项目主机信息

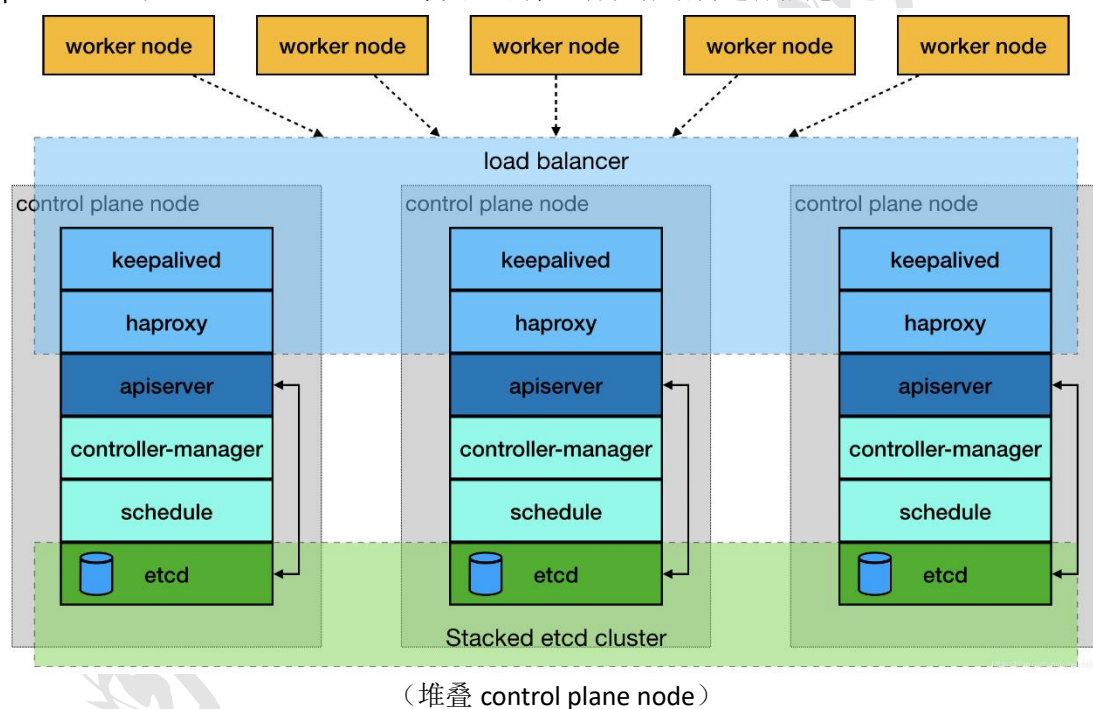
准备 6 台虚拟机，3 台 master 节点，3 台 node 节点，保证 master 节点数为 $\geq 3$  的奇数。  
硬件：2 核 CPU+、2G 内存+、硬盘 20G+

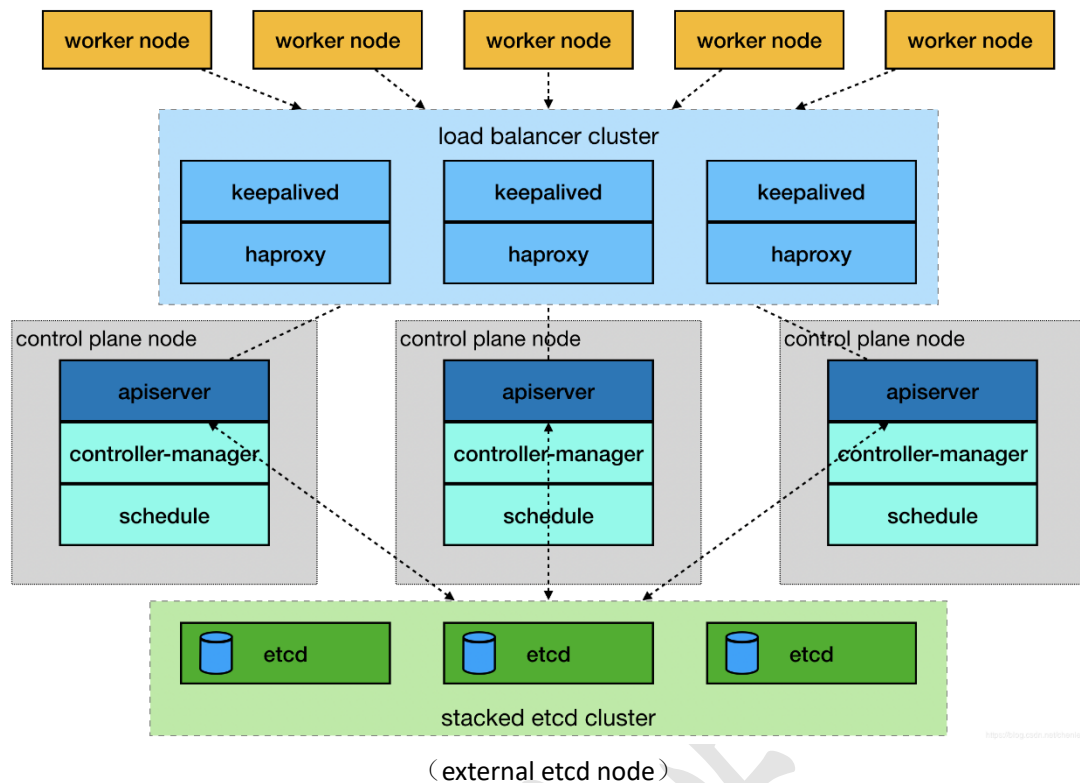
网络：所有机器网络互通、可以访问外网

操作系统	IP 地址	角色	主机名
CentOS7-x86-64	192.168.200.111	master	k8s-master1
CentOS7-x86-64	192.168.200.112	master	k8s-master2
CentOS7-x86-64	192.168.200.113	master	k8s-master3
CentOS7-x86-64	192.168.200.114	node	k8s-node1
CentOS7-x86-64	192.168.200.115	node	k8s-node2
CentOS7-x86-64	192.168.200.116	node	k8s-node3
	192.168.200.154	VIP	master.k8s.io

## 2.2、项目架构图

多 master 节点负载均衡的 kubernetes 集群。官网给出了两种拓扑结构：堆叠 control plane node 和 external etcd node，本文基于第一种拓扑结构进行搭建。





## 2.3、项目实施思路

master 节点需要部署 etcd、apiserver、controller-manager、scheduler 这 4 种服务，其中 etcd、controller-manager、scheduler 这三种服务 kubernetes 自身已经实现了高可用，在多 master 节点的情况下，每个 master 节点都会启动这三种服务，同一时间只有一个生效。因此要实现 kubernetes 的高可用，只需要 apiserver 服务高可用。

keepalived 是一种高性能的服务器高可用或热备解决方案，可以用来防止服务器单点故障导致服务中断的问题。keepalived 使用主备模式，至少需要两台服务器才能正常工作。比如 keepalived 将三台服务器搭建成一个集群，对外提供一个唯一 IP，正常情况下只有一台服务器上可以看到这个 IP 的虚拟网卡。如果这台服务异常，那么 keepalived 会立即将 IP 移动到剩下的两台服务器中的一台上，使得 IP 可以正常使用。

haproxy 是一款提供高可用性、负载均衡以及基于 TCP（第四层）和 HTTP（第七层）应用的代理软件，支持虚拟主机，它是免费、快速并且可靠的一种解决方案。使用 haproxy 负载均衡后端的 apiserver 服务，达到 apiserver 服务高可用的目的。

本文使用的 keepalived+haproxy 方案，使用 keepalived 对外提供稳定的入口，使用 haproxy 对内均衡负载。因为 haproxy 运行在 master 节点上，当 master 节点异常后，haproxy 服务也会停止，为了避免这种情况，我们在每一台 master 节点都部署 haproxy 服务，达到 haproxy 服务高可用的目的。由于多 master 节点会出现投票竞选的问题，因此 master 节点的数据最好是单数，避免票数相同的情况。

## 3、项目实施过程

### 3.1、系统初始化

关闭防火墙

```
[root@localhost ~]# systemctl stop firewalld
[root@localhost ~]# systemctl disable firewalld
```

关闭 selinux

```
[root@localhost ~]# sed -i 's/enforcing/disabled/' /etc/selinux/config
[root@localhost ~]# setenforce 0
```

关闭 swap

```
[root@localhost ~]# swapoff -a
[root@localhost ~]# sed -ri 's/.*swap.*#&/' /etc/fstab
```

修改主机名（根据主机角色不同，做相应修改）

```
hostname k8s-node3
bash
```

主机名映射

```
[root@k8s-master1 ~]# cat >> /etc/hosts << EOF
192.168.200.111 master1.k8s.io k8s-master1
192.168.200.112 master2.k8s.io k8s-master2
192.168.200.113 master3.k8s.io k8s-master3
192.168.200.114 node1.k8s.io k8s-node1
192.168.200.115 node2.k8s.io k8s-node2
192.168.200.116 node3.k8s.io k8s-node3
192.168.200.154 master.k8s.io k8s-vip
EOF
```

将桥接的 IPv4 流量传递到 iptables 的链

```
[root@k8s-master1 ~]# cat << EOF >> /etc/sysctl.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
[root@k8s-master1 ~]# modprobe br_netfilter
[root@k8s-master1 ~]# sysctl -p
```

时间同步

```
[root@k8s-master1 ~]# yum install ntpdate -y
```

```
[root@k8s-master1 ~]# ntpdate time.windows.com
```

## 3.2、配置部署 keepalived 服务

安装 Keepalived（所有 master 主机）

```
[root@k8s-master1 ~]# yum install -y keepalived
```

k8s-master1 节点配置

```
[root@k8s-master1 ~]# cat > /etc/keepalived/keepalived.conf <<EOF
! Configuration File for keepalived
global_defs {
    router_id k8s
}
vrrp_script check_haproxy {
    script "killall -0 haproxy"
    interval 3
    weight -2
    fall 10
    rise 2
}
vrrp_instance VI_1 {
    state MASTER
    interface ens33
    virtual_router_id 51
    priority 100
    advert_int 1
    authentication {
        auth_type PASS
        auth_pass 1111
    }
    virtual_ipaddress {
        192.168.200.154
    }
    track_script {
        check_haproxy
    }
}
EOF
```

k8s-master2 节点配置

```
[root@k8s-master2 ~]# cat > /etc/keepalived/keepalived.conf <<EOF
! Configuration File for keepalived
global_defs {
```

```
router_id k8s
}
vrrp_script check_haproxy {
    script "killall -0 haproxy"
    interval 3
    weight -2
    fall 10
    rise 2
}
vrrp_instance VI_1 {
    state BACKUP
    interface ens33
    virtual_router_id 51
    priority 90
    advert_int 1
    authentication {
        auth_type PASS
        auth_pass 1111
    }
    virtual_ipaddress {
        192.168.200.154
    }
    track_script {
        check_haproxy
    }
}
EOF
```

#### k8s-master3 节点配置

```
[root@k8s-master3 ~]# cat > /etc/keepalived/keepalived.conf <<EOF
! Configuration File for keepalived
global_defs {
    router_id k8s
}
vrrp_script check_haproxy {
    script "killall -0 haproxy"
    interval 3
    weight -2
    fall 10
    rise 2
}
vrrp_instance VI_1 {
    state BACKUP
    interface ens33
```



```

virtual_router_id 51
priority 80
adver_int 1
authentication {
    auth_type PASS
    auth_pass 1111
}
virtual_ipaddress {
    192.168.200.154
}
track_script {
    check_haproxy
}
}
EOF

```

启动和检查

所有 master 节点都要执行

```

[root@k8s-master1 ~]# systemctl start keepalived
[root@k8s-master1 ~]# systemctl enable keepalived

```

查看启动状态

```

[root@k8s-master1 ~]# systemctl status keepalived
Created symlink from /etc/systemd/system/multi-user.target.wants/keepalived.
service to /usr/lib/systemd/system/keepalived.service.[root@k8s-master1 ~]# systemctl status
keepalived
● keepalived.service - LVS and VRRP High Availability Monitor
   Loaded: loaded (/usr/lib/systemd/system/keepalived.service; enabled; vendor preset:
disabled)
   Active: active (running) since 日 2021-01-03 20:44:16 CST; 8s ago
 Main PID: 1670 (keepalived)
    CGroup: /system.slice/keepalived.service
            └─1670 /usr/sbin/keepalived -D
            └─1671 /usr/sbin/keepalived -D
            └─1672 /usr/sbin/keepalived -D

1 月 03 20:44:23 k8s-master1 Keepalived_vrrp[1672]: Sending gratuitous ARP on ens33 for
192.168.200.154
1 月 03 20:44:23 k8s-master1 Keepalived_vrrp[1672]: Sending gratuitous ARP on ens33 for
192.168.200.154
1 月 03 20:44:23 k8s-master1 Keepalived_vrrp[1672]: Sending gratuitous ARP on ens33 for
192.168.200.154
1 月 03 20:44:23 k8s-master1 Keepalived_vrrp[1672]: Sending gratuitous ARP on ens33 for
192.168.200.154
1 月 03 20:44:24 k8s-master1 Keepalived_vrrp[1672]: Sending gratuitous ARP on ens33 for

```

192.168.200.154

1 月 03 20:44:24 k8s-master1 Keepalived\_vrrp[1672]: VRRP\_Instance(VI\_1) Sending/queueing gratuitous ARPs on ens33 for 192.168.200.154

1 月 03 20:44:24 k8s-master1 Keepalived\_vrrp[1672]: Sending gratuitous ARP on ens33 for 192.168.200.154

1 月 03 20:44:24 k8s-master1 Keepalived\_vrrp[1672]: Sending gratuitous ARP on ens33 for 192.168.200.154

1 月 03 20:44:24 k8s-master1 Keepalived\_vrrp[1672]: Sending gratuitous ARP on ens33 for 192.168.200.154

1 月 03 20:44:24 k8s-master1 Keepalived\_vrrp[1672]: Sending gratuitous ARP on ens33 for 192.168.200.154

启动完成后在 master1 查看网络信息

```
[root@k8s-master1 ~]# ip a s ens33
```

```
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group
default qlen 1000
```

```
    link/ether 00:0c:29:ce:ef:5f brd ff:ff:ff:ff:ff:ff
```

```
    inet 192.168.200.111/24 brd 192.168.200.255 scope global ens33
```

```
        valid_lft forever preferred_lft forever
```

```
    inet 192.168.200.154/32 scope global ens33
```

```
        valid_lft forever preferred_lft forever
```

```
    inet6 fe80::20c:29ff:fece:ef5f/64 scope link
```

```
        valid_lft forever preferred_lft forever
```

### 3.3、配置部署 haproxy 服务

所有 master 主机安装 haproxy

```
[root@k8s-master1 ~]# yum install -y haproxy
```

每台 master 节点中的配置均相同，配置中声明了后端代理的每个 master 节点服务器，指定了 haproxy 的端口为 16443，因此 16443 端口为集群的入口。

```
[root@k8s-master1 ~]# cat > /etc/haproxy/haproxy.cfg << EOF
```

```
#-----
```

```
# Global settings
```

```
#-----
```

```
global
```

```
    log          127.0.0.1 local2
```

```
    chroot       /var/lib/haproxy
```

```
    pidfile      /var/run/haproxy.pid
```

```
    maxconn      4000
```

```
    user         haproxy
```

```
    group        haproxy
```

```
    daemon
```

```

stats socket /var/lib/haproxy/stats
#-----
# common defaults that all the 'listen' and 'backend' sections will
# use if not designated in their block
#-----
defaults
    mode                http
    log                 global
    option              httplog
    option              dontlognull
    option http-server-close
    option forwardfor   except 127.0.0.0/8
    option              redispatch
    retries             3
    timeout http-request 10s
    timeout queue       1m
    timeout connect     10s
    timeout client      1m
    timeout server      1m
    timeout http-keep-alive 10s
    timeout check       10s
    maxconn             3000
#-----
# kubernetes apiserver frontend which proxys to the backends
#-----
frontend kubernetes-apiserver
    mode                tcp
    bind                *:16443
    option              tcplog
    default_backend     kubernetes-apiserver
#-----
#round robin balancing between the various backends
#-----
backend kubernetes-apiserver
    mode                tcp
    balance              roundrobin
    server              master1.k8s.io    192.168.200.111:6443 check
    server              master2.k8s.io    192.168.200.112:6443 check
    server              master3.k8s.io    192.168.200.113:6443 check
#-----
# collection haproxy statistics message
#-----
listen stats
    bind                *:1080

```

```
stats auth      admin:awesomePassword
stats refresh   5s
stats realm     HAProxy\ Statistics
stats uri       /admin?stats
EOF
```

启动和检查

所有 master 节点都要执行

```
[root@k8s-master1 ~]# systemctl start haproxy
[root@k8s-master1 ~]# systemctl enable haproxy
```

查看启动状态

```
[root@k8s-master1 ~]# systemctl status haproxy
[root@k8s-master1 ~]# systemctl status haproxy
● haproxy.service - HAProxy Load Balancer
   Loaded: loaded (/usr/lib/systemd/system/haproxy.service; enabled; vendor preset: disabled)
   Active: active (running) since 日 2021-01-03 20:45:39 CST; 11s ago
   Main PID: 1778 (haproxy-systemd)
   CGroup: /system.slice/haproxy.service
           └─1778 /usr/sbin/haproxy-systemd-wrapper -f /etc/haproxy/haproxy.cfg -p /run/haproxy.pid
           └─1779 /usr/sbin/haproxy -f /etc/haproxy/haproxy.cfg -p /run/haproxy.pid -Ds
           └─1780 /usr/sbin/haproxy -f /etc/haproxy/haproxy.cfg -p /run/haproxy.pid -Ds

1月 03 20:45:39 k8s-master1 systemd[1]: Started HAProxy Load Balancer.
1月 03 20:45:39 k8s-master1 haproxy-systemd-wrapper[1778]: haproxy-systemd-wrapper: executing /usr/sbin/haproxy -f /etc/haproxy/haproxy.cfg -p /ru...pid -Ds
1月 03 20:45:39 k8s-master1 haproxy-systemd-wrapper[1778]: [WARNING] 002/204539 (1779) : config : 'option forwardfor' ignored for frontend 'kubern...P mode.
1月 03 20:45:39 k8s-master1 haproxy-systemd-wrapper[1778]: [WARNING] 002/204539 (1779) : config : 'option forwardfor' ignored for backend 'kuberne...P mode.
Hint: Some lines were ellipsized, use -l to show in full.
```

检查端口

```
[root@k8s-master1 ~]# netstat -lntup | grep haproxy
tcp        0      0 0.0.0.0:1080          0.0.0.0:*            LISTEN
1780/haproxy
tcp        0      0 0.0.0.0:16443         0.0.0.0:*            LISTEN
1780/haproxy
udp        0      0 0.0.0.0:36433        0.0.0.0:*
1779/haproxy
```

### 3.4、配置部署 Docker 服务

所有主机上分别部署 Docker 环境，因为 Kubernetes 对容器的编排需要 Docker 的支持。

```
[root@k8s-master ~]# wget -O /etc/yum.repos.d/CentOS-Base.repo
http://mirrors.aliyun.com/repo/Centos-7.repo
[root@k8s-master ~]# yum install -y yum-utils device-mapper-persistent-data lvm2
```

使用 YUM 方式安装 Docker 时，推荐使用阿里的 YUM 源。

```
[root@k8s-master ~]# yum-config-manager --add-repo https://mirrors.aliyun.com/docker-
ce/linux/centos/docker-ce.repo
[root@k8s-master ~]# yum clean all && yum makecache fast

[root@k8s-master ~]# yum -y install docker-ce
[root@k8s-master ~]# systemctl start docker
[root@k8s-master ~]# systemctl enable docker
```

镜像加速器（所有主机配置）

```
[root@k8s-master ~]# cat << END > /etc/docker/daemon.json
{
    "registry-mirrors": [ "https://nyakyfun.mirror.aliyuncs.com" ]
}
END
[root@k8s-master ~]# systemctl daemon-reload
[root@k8s-master ~]# systemctl restart docker
```

### 3.5、部署 kubelet kubeadm kubectl 工具

使用 YUM 方式安装 Kubernetes 时，推荐使用阿里的 yum。

所有主机配置

```
[root@k8s-master ~]# cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://mirrors.aliyun.com/kubernetes/yum/repos/kubernetes-el7-x86_64/
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://mirrors.aliyun.com/kubernetes/yum/doc/yum-key.gpg
        https://mirrors.aliyun.com/kubernetes/yum/doc/rpm-package-key.gpg
EOF

[root@k8s-master ~]# ls /etc/yum.repos.d/
backup  Centos-7.repo  CentOS-Media.repo  CentOS-x86_64-kernel.repo  docker-ce.repo
```

```
kubernetes.repo
```

安装 kubelet kubeadm kubectl

所有主机配置

```
[root@k8s-master ~]# yum install -y kubelet kubeadm kubectl  
[root@k8s-master ~]# systemctl enable kubelet
```

## 3.6、部署 Kubernetes Master

在具有 vip 的 master 上操作。此处的 vip 节点为 k8s-master1。

创建 kubeadm-config.yaml 文件

```
[root@k8s-master1 ~]# cat > kubeadm-config.yaml << EOF
```

```
apiServer:
```

```
  certSANs:
```

```
    - k8s-master1  
    - k8s-master2  
    - k8s-master3  
    - master.k8s.io  
    - 192.168.200.111  
    - 192.168.200.112  
    - 192.168.200.113  
    - 192.168.200.154  
    - 127.0.0.1
```

```
  extraArgs:
```

```
    authorization-mode: Node,RBAC
```

```
  timeoutForControlPlane: 4m0s
```

```
apiVersion: kubeadm.k8s.io/v1beta1
```

```
certificatesDir: /etc/kubernetes/pki
```

```
clusterName: kubernetes
```

```
controlPlaneEndpoint: "master.k8s.io:6443"
```

```
controllerManager: {}
```

```
dns:
```

```
  type: CoreDNS
```

```
etcd:
```

```
  local:
```

```
    dataDir: /var/lib/etcd
```

```
imageRepository: registry.aliyuncs.com/google_containers
```

```
kind: ClusterConfiguration
```

```
kubernetesVersion: v1.20.0
```

```
networking:
```

```
  dnsDomain: cluster.local
```

```
  podSubnet: 10.244.0.0/16
```

```
  serviceSubnet: 10.1.0.0/16
```

```
scheduler: {}
EOF
```

查看所需镜像信息

```
[root@k8s-master1 ~]# kubeadm config images list --config kubeadm-config.yaml
W0103 21:00:29.765926      3145 common.go:77] your configuration file uses a deprecated
API spec: "kubeadm.k8s.io/v1beta1". Please use 'kubeadm config migrate
--old-config old.yaml --new-config new.yaml', which will write the new, similar spec using a
newer API version.
registry.aliyuncs.com/google_containers/kube-apiserver:v1.20.0
registry.aliyuncs.com/google_containers/kube-controller-manager:v1.20.0
registry.aliyuncs.com/google_containers/kube-scheduler:v1.20.0
registry.aliyuncs.com/google_containers/kube-proxy:v1.20.0
registry.aliyuncs.com/google_containers/pause:3.2
registry.aliyuncs.com/google_containers/etcd:3.4.13-0
registry.aliyuncs.com/google_containers/coredns:1.7.0
```

上传 k8s 所需的镜像并导入（所有 master 主机）

```
[root@k8s-master1 ~]# cd master/
[root@k8s-master1 master]# ls
coredns_1.7.0.tar kube-apiserver_v1.20.0.tar kube-proxy_v1.20.0.tar
pause_3.2.tar etcd_3.4.13-0.tar kube-controller-manager_v1.20.0.tar kube-
scheduler_v1.20.0.tar
[root@k8s-master1 master]# ls | while read line
do
docker load < $line
done
```

使用 kubeadm 命令初始化 k8s

```
[root@k8s-master1 ~]# kubeadm init --config kubeadm-config.yaml

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

You can now join any number of control-plane nodes by copying certificate authorities
and service account keys on each node and then running the following as root:

kubeadm join master.k8s.io:6443 --token vhfqzv.c398m3ohhpl8avd0 \ 加入master时使用
--discovery-token-ca-cert-hash sha256:40fc58440b1ea96c64816050c197b6ef0cd8ac42f687d0e54b8b2f9c392a5617 \
--control-plane

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join master.k8s.io:6443 --token vhfqzv.c398m3ohhpl8avd0 \ 加入node时使用
--discovery-token-ca-cert-hash sha256:40fc58440b1ea96c64816050c197b6ef0cd8ac42f687d0e54b8b2f9c392a5617
```

初始化中的错误：

```
[ERROR FileContent--proc-sys-net-bridge-bridge-nf-call-iptables]: /proc/sys/net/bridge/bridge-nf-call-iptables contents are not set to 1
```

执行以下命令后重新执行初始化命令

```
echo "1" >/proc/sys/net/bridge/bridge-nf-call-iptables
```

根据初始化的结果操作

```
[root@k8s-master1 ~]# mkdir -p $HOME/.kube
[root@k8s-master1 ~]# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
[root@k8s-master1 ~]# sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

查看集群状态

```
[root@k8s-master1 manifests]# kubectl get cs
```

NAME	STATUS	MESSAGE	ERROR
controller-manager	Unhealthy	Get "http://127.0.0.1:10252/healthz": dial tcp 127.0.0.1:10252: connect: connection refused	ERROR
scheduler	Unhealthy	Get "http://127.0.0.1:10251/healthz": dial tcp 127.0.0.1:10251: connect: connection refused	
etcd-0	Healthy	{"health": "true"}	

注意：出现以上错误情况，是因为/etc/kubernetes/manifests/下的 kube-controller-manager.yaml 和 kube-scheduler.yaml 设置的默认端口为 0 导致的，解决方式是注释掉对应的 port 即可

修改 kube-controller-manager.yaml 文件

```
[root@k8s-master1 ~]# cat /etc/kubernetes/manifests/kube-controller-manager.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    component: kube-controller-manager
    tier: control-plane
  name: kube-controller-manager
  namespace: kube-system
spec:
  containers:
    - command:
      - kube-controller-manager
      - --allocate-node-cidrs=true
      - --authentication-kubeconfig=/etc/kubernetes/controller-manager.conf
      - --authorization-kubeconfig=/etc/kubernetes/controller-manager.conf
      - --bind-address=127.0.0.1
      - --client-ca-file=/etc/kubernetes/pki/ca.crt
      - --cluster-cidr=10.244.0.0/16
      - --cluster-name=kubernetes
      - --cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt
      - --cluster-signing-key-file=/etc/kubernetes/pki/ca.key
      - --controllers=*,bootstrapsigner,tokencleaner
      - --kubeconfig=/etc/kubernetes/controller-manager.conf
      - --leader-elect=true
    #   - --port=0
      - --requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt
```

修改 kube-scheduler.yaml 文件



```
[root@k8s-master1 ~]# cat /etc/kubernetes/manifests/kube-scheduler.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    component: kube-scheduler
    tier: control-plane
  name: kube-scheduler
  namespace: kube-system
spec:
  containers:
  - command:
    - kube-scheduler
    - --authentication-kubeconfig=/etc/kubernetes/scheduler.conf
    - --authorization-kubeconfig=/etc/kubernetes/scheduler.conf
    - --bind-address=127.0.0.1
    - --kubeconfig=/etc/kubernetes/scheduler.conf
    - --leader-elect=true
    #   --port=0
    image: registry.aliyuncs.com/google_containers/kube-scheduler:v1.20.0
```

查看集群状态

```
[root@k8s-master1 ~]# kubectl get cs
Warning: v1 ComponentStatus is deprecated in v1.19+
NAME                STATUS    MESSAGE                       ERROR
controller-manager  Healthy   ok
scheduler            Healthy   ok
etcd-0               Healthy   {"health":"true"}
```

查看 pod 信息

```
[root@k8s-master1 ~]# kubectl get pods -n kube-system
NAME                                READY   STATUS    RESTARTS   AGE
coredns-7f89b7bc75-52kjr           0/1     Pending   0           8m30s
coredns-7f89b7bc75-7h8js           0/1     Pending   0           8m30s
etcd-k8s-master1                    1/1     Running   0           8m38s
kube-apiserver-k8s-master1          1/1     Running   0           8m38s
kube-controller-manager-k8s-master1 1/1     Running   0           2m11s
kube-proxy-s24dn                    1/1     Running   0           8m31s
kube-scheduler-k8s-master1          1/1     Running   0           119s
```

查看节点信息

```
[root@k8s-master1 ~]# kubectl get nodes
NAME                STATUS    ROLES                  AGE   VERSION
k8s-master1        NotReady  control-plane,master   9m24s v1.20.0
```

### 3.7、安装集群网络

在 k8s-master1 节点执行

```
[root@k8s-master1 ~]# docker load < flannel_v0.12.0-amd64.tar
```

```
[root@k8s-master1 ~]# kubectl apply -f kube-flannel.yml
```

再次查看节点信息：

```
[root@k8s-master1 ~]# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8s-master1	Ready	control-plane,master	8m13s	v1.20.1

### 3.8、添加 master 节点

在 k8s-master2 和 k8s-master3 节点创建文件夹

```
[root@k8s-master2 ~]# mkdir -p /etc/kubernetes/pki/etcd
```

```
[root@k8s-master3 ~]# mkdir -p /etc/kubernetes/pki/etcd
```

在 k8s-master1 节点执行

从 k8s-master1 复制秘钥和相关文件到 k8s-master2 和 k8s-master3

```
[root@k8s-master1 ~]# scp /etc/kubernetes/admin.conf
root@192.168.200.112:/etc/kubernetes
[root@k8s-master1 ~]# scp /etc/kubernetes/admin.conf
root@192.168.200.113:/etc/kubernetes

[root@k8s-master1 ~]# scp /etc/kubernetes/pki/{ca.*,sa.*,front-proxy-ca.*}
root@192.168.200.112:/etc/kubernetes/pki
[root@k8s-master1 ~]# scp /etc/kubernetes/pki/{ca.*,sa.*,front-proxy-ca.*}
root@192.168.200.113:/etc/kubernetes/pki

[root@k8s-master1 ~]# scp /etc/kubernetes/pki/etcd/ca.*
root@192.168.200.112:/etc/kubernetes/pki/etcd
[root@k8s-master1 ~]# scp /etc/kubernetes/pki/etcd/ca.*
root@192.168.200.113:/etc/kubernetes/pki/etcd
```

将其他 master 节点加入集群

注意：kubeadm init 生成的 token 有效期只有 1 天，生成不过期 token

```
[root@k8s-master1 manifests]# kubeadm token create --ttl 0 --print-join-command
[root@k8s-master1 manifests]# kubeadm token list
```

```
[root@k8s-master1 ~]# kubeadm token create --ttl 0 --print-join-command
kubeadm join master.k8s.io:6443 --token 5zvcjk.i6r9vjc869svkke --discovery-token-ca-cert-hash sha256:40fc58440b1ea96c64816050c197b6ef0cd8ac42f687d0e54b8b2f9c392a5617
[root@k8s-master1 ~]# kubeadm token list
```

TOKEN	TTL	EXPIRES	USAGES	DESCRIPTION	EXTRA GROUPS
5zvcjk.i6r9vjc869svkke	<forever>	<never>	authentication,signing	<none>	system:bootstrappers:kube
adm:default-node-token					
vhqazv.c398m3ohhp18avd0	23h	2021-01-04T21:06:38+08:00	authentication,signing	<none>	system:
bootstrappers:kubeadm:default-node-token					

k8s-master2 和 k8s-master3 都需要加入

```
[root@k8s-master2 ~]# kubeadm join master.k8s.io:6443 --token 5zvcjk.i6r9vjc869svkke \
--discovery-token-ca-cert-hash
```

```
sha256:40fc58440b1ea96c64816050c197b6ef0cd8ac42f687d0e54b8b2f9c392a5617 \
--control-plane

[root@k8s-master2 ~]# mkdir -p $HOME/.kube
[root@k8s-master2 ~]# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
[root@k8s-master2 ~]# sudo chown $(id -u):$(id -g) $HOME/.kube/config

[root@k8s-master2 ~]# docker load < flannel_v0.12.0-amd64.tar
```

```
[root@k8s-master1 ~]# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8s-master1	Ready	control-plane,master	19m	v1.20.0
k8s-master2	Ready	control-plane,master	2m20s	v1.20.0
k8s-master3	Ready	control-plane,master	95s	v1.20.0

```
[root@k8s-master1 manifests]# kubectl get pods --all-namespaces
```

```
[root@k8s-master1 ~]# kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-7f89b7bc75-52kjr	1/1	Running	0	19m
kube-system	coredns-7f89b7bc75-7h8js	1/1	Running	0	19m
kube-system	etcd-k8s-master1	1/1	Running	0	19m
kube-system	etcd-k8s-master2	1/1	Running	0	2m34s
kube-system	etcd-k8s-master3	1/1	Running	0	43s
kube-system	kube-apiserver-k8s-master1	1/1	Running	0	19m
kube-system	kube-apiserver-k8s-master2	1/1	Running	0	2m38s
kube-system	kube-apiserver-k8s-master3	1/1	Running	1	110s
kube-system	kube-controller-manager-k8s-master1	1/1	Running	1	13m
kube-system	kube-controller-manager-k8s-master2	1/1	Running	0	2m38s
kube-system	kube-controller-manager-k8s-master3	1/1	Running	0	36s
kube-system	kube-flannel-ds-amd64-nfxcc	1/1	Running	0	115s
kube-system	kube-flannel-ds-amd64-pb94t	1/1	Running	0	2m40s
kube-system	kube-flannel-ds-amd64-x8lck	1/1	Running	0	9m31s
kube-system	kube-proxy-gdb5l	1/1	Running	0	2m40s
kube-system	kube-proxy-s24dn	1/1	Running	0	19m
kube-system	kube-proxy-x89vb	1/1	Running	0	115s
kube-system	kube-scheduler-k8s-master1	1/1	Running	1	12m
kube-system	kube-scheduler-k8s-master2	1/1	Running	0	2m38s
kube-system	kube-scheduler-k8s-master3	1/1	Running	0	35s

### 3.9、加入 Kubernetes Node

直接在 node 节点服务器上执行 k8s-master1 初始化成功后的消息即可：

```
[root@k8s-node1 ~]# kubeadm join master.k8s.io:6443 --token 5zvcjk.i6r9vijk869svkke \
--discovery-token-ca-cert-hash
sha256:40fc58440b1ea96c64816050c197b6ef0cd8ac42f687d0e54b8b2f9c392a5617

[root@k8s-node1 ~]# docker load < flannel_v0.12.0-amd64.tar
```

查看节点信息

```
[root@k8s-master1 ~]# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8s-master1	Ready	control-plane,master	23m	v1.20.0

k8s-master2	Ready	control-plane,master	6m42s	v1.20.0
k8s-master3	Ready	control-plane,master	5m57s	v1.20.0
k8s-node1	Ready	<none>	82s	v1.20.0
k8s-node2	Ready	<none>	77s	v1.20.0
k8s-node3	Ready	<none>	60s	v1.20.0

## 3.10、测试 Kubernetes 集群

所有 node 主机导入测试镜像

```
[root@k8s-node1 ~]# docker load < nginx-1.19.tar
[root@k8s-node1 ~]# docker tag nginx nginx:1.19.6
```

在 Kubernetes 集群中创建一个 pod，验证是否正常运行。

```
[root@k8s-master1 ~]# mkdir demo
[root@k8s-master1 ~]# cd demo
[root@k8s-master1 demo]# vim nginx-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.19.6
          ports:
            - containerPort: 80
```

创建完 Deployment 的资源清单之后，使用 create 执行资源清单来创建容器。通过 get pods 可以查看到 Pod 容器资源已经自动创建完成。

```
[root@k8s-master1 demo]# kubectl create -f nginx-deployment.yaml
deployment.apps/nginx-deployment created
```

```
[root@k8s-master1 demo]# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-76ccf9dd9d-c5bjw	1/1	Running	0	10s
nginx-deployment-76ccf9dd9d-jxm9z	1/1	Running	0	10s
nginx-deployment-76ccf9dd9d-kzdfz	1/1	Running	0	10s

```
[root@k8s-master1 ~]# kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP
nginx-deployment-76ccf9dd9d-c5bjw	1/1	Running	0	62m	10.244.4.2
k8s-node2	<none>	<none>			
nginx-deployment-76ccf9dd9d-jxm9z	1/1	Running	0	62m	10.244.5.2
k8s-node3	<none>	<none>			
nginx-deployment-76ccf9dd9d-kzdfz	1/1	Running	0	62m	10.244.3.2
k8s-node1	<none>	<none>			

### 创建 Service 资源清单

在创建的 nginx-service 资源清单中，定义名称为 nginx-service 的 Service、标签选择器为 app:nginx、type 为 NodePort 指明外部流量可以访问内部容器。在 ports 中定义暴露的端口库号列表，对外暴露访问的端口是 80，容器内部的端口也是 80。

```
[root@k8s-master1 demo]# vim nginx-service.yaml
```

```
kind: Service
apiVersion: v1
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  type: NodePort
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

```
[root@k8s-master1 demo]# kubectl create -f nginx-service.yaml
```

```
service/nginx-service created
```

```
[root@k8s-master1 demo]# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.1.0.1	<none>	443/TCP	38m
nginx-service	NodePort	10.1.161.204	<none>	80:30373/TCP	4s

通过浏览器访问 nginx: <http://master.k8s.io:30373> 域名或者 VIP 地址

```
[root@k8s-master1 demo]# elinks --dump http://master.k8s.io:30373
```

```
Welcome to nginx!
```

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [1]nginx.org.  
Commercial support is available at [2]nginx.com.

Thank you for using nginx.

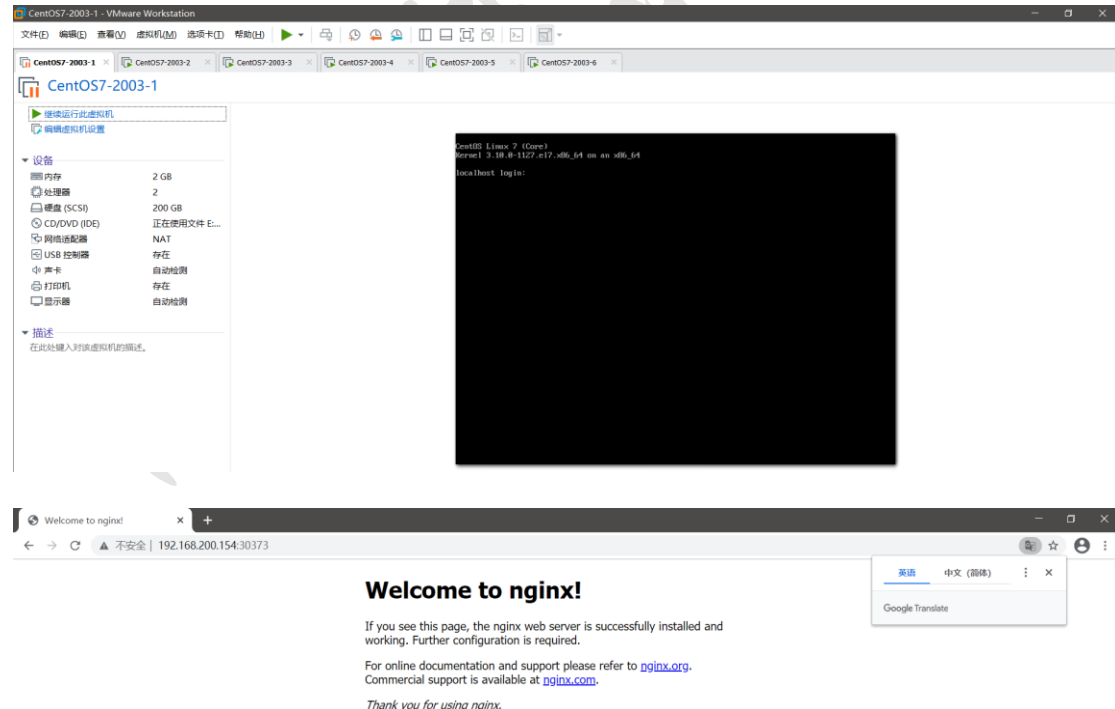
## References

### Visible links

1. <http://nginx.org/>
2. <http://nginx.com/>



挂起 k8s-master1 节点，刷新页面还是能访问 nginx，说明高可用集群部署成功。



检查会发现 VIP 已经转移到 k8s-master2 节点上

```
[root@k8s-master2 ~]# ip a s ens33
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group
default qlen 1000
```

```
link/ether 00:0c:29:3a:76:54 brd ff:ff:ff:ff:ff:ff
inet 192.168.200.112/24 brd 192.168.200.255 scope global ens33
    valid_lft forever preferred_lft forever
inet 192.168.200.154/32 scope global ens33
    valid_lft forever preferred_lft forever
inet6 fe80::20c:29ff:fe3a:7654/64 scope link
    valid_lft forever preferred_lft forever
```

至此 Kubernetes 企业级高可用环境完美实现。

## 4、项目总结

- 1、集群中只要有一个 master 节点正常运行就可以正常对外提供业务服务。
- 2、如果需要在 master 节点使用 kubectl 相关的命令，必须保证至少有 2 个 master 节点正常运行才可以使用，不然会有 Unable to connect to the server: net/http: TLS handshake timeout 这样的错误。
- 3、节点故障时 pod 自动转移：当 pod 所在的节点宕机后，根据 controller-manager 的 pod-eviction-timeout 配置，默认是 5 分钟，5 分钟后 k8s 会把 pod 状态设置为 unknown，然后在其它节点启动 pod。当故障节点恢复后，k8s 会删除故障节点上面的 unknown pod。如果你想立即强制迁移，可以用 kubectl drain nodename
- 4、为了保证集群的高可用性，建议 master 节点和 node 节点至少分别部署 3 台及以上，且 master 节点应该部署基数个实例(3、5、7、9)。